

# Cardinality Estimation of Distinct Items - An Overview

Sebastian Balsam

November 18, 2025

## Abstract

In this paper I want to give an overview of different solutions for the Count Distinct Problem...

## 1 Introduction

The cardinality estimation problem or count-distinct problem is about finding the number of distinct elements in a data stream with repeated elements. These elements could be url's, unique users, sensor data or IP addresses passing through a data stream. A simple solution would be to use a list and add an item to this list each time we encounter an unseen item. With this approach we can give an exact answer to the query, how many distinct elements have been seen. Unfortunately if millions of distinct elements are present in a data stream, with this approach, we will soon hit storage boundaries and the performance will deteriorate.

As we often do not need an exact answer, streaming algorithms have been developed, that give an approximation that is mostly good enough and bound to a fixed storage size. There different approaches to solve this problem. Sampling techniques are used as an estimate by sampling a subset of items in the set and use the subset as estimator, e.g the CVM algorithm by Chakraborty et al. [Chakraborty et al., 2023]. While sampling methods generally give a good estimate, they can fail in certain situations. If rare elements are not in the sampled set, they would not be counted and we would underestimate. For database optimization, machine learning techniques have been used recently for cardinality estimation [Liu et al., 2015, Woltmann et al., 2019, Schwabe and Acosta, 2024].

The technique I want to focus on in this paper are sketch techniques that implement a certain data structure and an accompanying algorithm to store information efficient for cardinality estimation.

## 2 Probabilistic Counting with Stochastic Averaging

- Overview of PCSA - Problem description

### 2.1 Flajolet and Martin

Before the advent of the internet there were not many data stream applications as we have today. But databases existed and began to grow in size. As table sizes grew, evaluation strategies became important, how to handle such big data

tables for join operations. Query optimizers were developed that could find the optimal strategy.

In this context Flajolet and Martin developed in 1985 a method to estimate the number of distinct items in a set in a space efficient way[Flajolet and Martin, 1985]. The idea behind the method is to use a hash function that maps  $n$  items uniformly to integers. If we look at the binary representation of these integers, and check the longest runs of zeros from the right, in about  $n/2$  cases, we have a '1' already in the first bit from the right. In about  $n/4$  cases, we have a '1' in the second last bit and so on. That means, based on the number of the rightmost '1', we can estimate the number of items. When we add an item, we use the algorithm as described in Algorithm 1 to add set the position of the rightmost '1' in the bitmap.

Flajolet and Martin found that they get better result, if multiple bins of bitmaps are used. With a higher number of bins (the value of  $nmap$ ) we use for the average, the estimation quality raises, but at the same time we need more memory to store the bitmap. They use the hash value to decide the bin (by using the modulo operation) and save the rightmost '1' of the binary representation in this bin.

---

**Algorithm 1** Adding an item in the Flajolet-Martin algorithm.

```
1: function  $\rho(val)$ 
2:   return position of the first 1 bit in val.
3: end function
4: function ADDITEM( $x$ )
5:    $hashedx \leftarrow hash(x)$ 
6:    $\alpha \leftarrow hashedx \bmod nmap$ 
7:    $index \leftarrow \rho(hashedx \bmod nmap)$ 
8:    $BITMAP[\alpha][index] = 1$ 
9: end function
```

---

When an estimate is needed, an average of the leftmost zeros are used, as shown in algorithm 2. The magic number  $\rho = 0.77351$  is used, as the expected Value of R - the position of leftmost zero in the bitmap is

$$E(R) \approx \log_2 \rho n$$

- Compressed (Scheuermann)

## 3 HyperLogLog, 2007

A first refinement of the Flajolet & Martin algorithm came in 2003 by Durand and Flajolet [Durand and Flajolet, 2003].

---

**Algorithm 2** Get an estimate in the Flajolet-Martin algorithm.

---

```

1: function QUERY
2:    $\rho \leftarrow 0.77351$ 
3:   max  $\leftarrow 32$                                  $\triangleright$  32 bit per bin
4:   for  $i := 0$  to  $nmap - 1$  do
5:     while  $BITMAP[i][R] == 1$  and  $R < max$  do
6:        $R \leftarrow R + 1$ 
7:        $S \leftarrow S + R$ 
8:     end while
9:   end for
10:  return  $\text{int}(nmap/\rho * 2^{S/nmap})$ 
11: end function

```

---

Later in 2007 it was further refined [Flajolet et al., 2007] to improve the accuracy even further.

The idea is quite simple. When we try to estimate the number of items, in the FM85 algorithm, we look for the first '0' in the bitmap. All the other '1' before and any bits after this position are not used. So instead of saving the complete bitmap, it should be enough to just save the maximum number of ones from the left. In addition since we have a maximum of 32 bits originally in the bitmap, we only need 5 bits ( $2^5 = 32$ ) for this number to be stored for each bin. This is an additional memory reduction.

## 4 HyperLogLog++, 2013

## 5 LogLogBeta, 2016

## 6 ExtendedHyperLogLog, 2023

<https://arxiv.org/pdf/2106.06525.pdf>

## 7 HLL-TailCut, 2023

[https://topoer-seu.github.io/PersonalPage/csqjxiao\\_files/papers/INFOCOM2023.pdf](https://topoer-seu.github.io/PersonalPage/csqjxiao_files/papers/INFOCOM2023.pdf)

## 8 Conclusion

Table 1: Table showing the different methods...

Algorithm	Memory	Comment
Flajolet & Martin	23	clever
HLL	45	here
HLL++	23	23
ExHLL	35	fsd
HLL-TailCut	23	sdf

## References

- [Chakraborty et al., 2023] Chakraborty, Vinodchandran, and Meel (2023). Distinct elements in streams: An algorithm for the (text) book.
- [Durand and Flajolet, 2003] Durand and Flajolet (2003). Loglog counting of large cardinalities.
- [Flajolet et al., 2007] Flajolet, Fusy, Gandouet, and Meunier (2007). Hyperloglog: the analysis of a near-optimal cardinality estimation algorithm.
- [Flajolet and Martin, 1985] Flajolet, P. and Martin, G. N. (1985). Probabilistic counting algorithms for data base applications.
- [Liu et al., 2015] Liu, Xu, Yu, Covinelli, and ZuZarte (2015). Cardinality estimation using neural networks.
- [Schwabe and Acosta, 2024] Schwabe and Acosta (2024). Cardinality estimation over knowledge graphs with embeddings and graph neural networks.
- [Woltmann et al., 2019] Woltmann, Hartmann, Thiele, Habich, and Lehner (2019). Cardinality estimation with INFOCOMP models.